

PD6: Prediction of the Release Year of a Song

MTI850 - Big Data Analytics

Fall 2025

Fixed for mllib to ml

Équipe 9

|



This assignment covers a common supervised learning pipeline, using a subset of the [Million Song Dataset](#) from the [UCI Machine Learning Repository](#). Our goal is to train a linear regression model to predict the release year of a song given a set of audio features. The subset has 6,724 samples out of 515,345 in the original dataset.

This assignment will cover:

- *Part 1:* Read and parse the initial dataset
 - *Visualization 1:* Features
 - *Visualization 2:* Shifting labels
- *Part 2:* Create and evaluate a baseline model
 - *Visualization 3:* Predicted vs. actual
- *Part 3:* Train (via gradient descent) and evaluate a linear regression model
 - *Visualization 4:* Training error
- *Part 4:* Train using SparkML library and tune hyperparameters via grid search
 - *Visualization 5:* Best model's predictions
 - *Visualization 6:* Hyperparameter heat map
- *Part 5:* Add interactions between features

Note that, for reference, you can look up the details of:

- the relevant Spark methods in [Spark's RDD Python API](#) and [Spark's DataFrame Python API](#)
- the relevant NumPy methods in the [NumPy Reference](#)

Part 0: Preparing the environment

```
In [1]: import os
import findspark
findspark.init()

# Test module for MTI850
import testmti850

# Util module for MTI850
import utilmti850

import pyspark
from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .master("local") \
    .appName("Music Year Prediction") \
    .getOrCreate()
```

```
WARNING: Using incubator modules: jdk.incubator.vector
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
25/11/28 00:04:56 WARN Utils: Your hostname, mti850vm, resolves to a loopback address: 127.0.1.1; using 10.0.2.15 instead (on interface enp0s8)
25/11/28 00:04:56 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to an other address
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
25/11/28 00:04:56 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

Part 1: Read and parse the initial dataset

(1a) Download, load, and check the data

The raw data is currently stored in a text file. Let's download it from Moodle and make it available in the hdfs filesystem.

```
start-dfs.sh && start-yarn.sh
hdfs dfs -put millionsong.txt /millionsong.txt
```

Then, we will store this raw data in as a DataFrame, with each element of the DataFrame representing a data point as a comma-delimited string. Each string starts with the label (a year) followed by numerical audio features. Use the DataFrame [count method](#) to

check how many data points we have. Then use the `take` method to create and print out a list of the first 5 data points in their initial string format.

```
In [3]: raw_data_df = spark.read.text("./millionsong.txt")
raw_data_df.show(5)
```

```
+-----+
|      value|
+-----+
|2001.0,0.88412373...|
|2001.0,0.85441194...|
|2001.0,0.90898297...|
|2001.0,0.84252521...|
|2001.0,0.90930328...|
+-----+
only showing top 5 rows
```

```
In [6]: # TODO: Replace <FILL IN> with appropriate code
num_points = raw_data_df.count()
print(num_points)

sample_points = raw_data_df.take(5)
print(sample_points)
```

```
6724
[Row(value='2001.0,0.884123733793,0.610454259079,0.600498416968,0.47466921
2493,0.247232680947,0.357306088914,0.344136412234,0.339641227335,0.6008588
40135,0.425704689024,0.60491501652,0.419193351817'), Row(value='2001.0,0.8
54411946129,0.604124786151,0.593634078776,0.495885413963,0.26630783093
6,0.261472105188,0.506387076327,0.464453565511,0.665798573683,0.5429689887
66,0.58044428577,0.445219373624'), Row(value='2001.0,0.908982970575,0.6320
63159227,0.557428975183,0.498263761394,0.276396052336,0.312809861625,0.448
530069406,0.448674249968,0.649791323916,0.489868662682,0.591908113534,0.45
00023818'), Row(value='2001.0,0.842525219898,0.561826888508,0.50871525969
2,0.443531142139,0.296733836002,0.250213568176,0.488540873206,0.3605087476
59,0.575435243185,0.361005878554,0.678378718617,0.409036786173'), Row(valu
e='2001.0,0.909303285534,0.653607720915,0.585580794716,0.473250503005,0.25
1417011835,0.326976795524,0.40432273022,0.371154511756,0.629401917965,0.48
2243251755,0.566901413923,0.463373691946')]
```

```
In [7]: # TEST Load and check the data (1a)
testmti850.Test.assertEquals(num_points, 6724, 'incorrect value for num_p
testmti850.Test.assertEquals(len(sample_points), 5, 'incorrect length for
```

1 test passed.
1 test passed.

(1b) Using LabeledPoint

In MLlib, labeled training instances are stored using the `LabeledPoint` object. Write the `parse_points` function that takes, as input, a DataFrame of comma-separated strings. We'll pass it the `raw_data_df` DataFrame.

It should parse each row in the DataFrame into individual elements, using Spark's `select` and `split` methods.

For example, split

```
"2001.0,0.884,0.610,0.600,0.474,0.247,0.357,0.344,0.33,0.600,0.425,  
into ['2001.0', '0.884', '0.610', '0.600', '0.474', '0.247',  
'0.357', '0.344', '0.33', '0.600', '0.425', '0.60', '0.419'] .
```

The first value in the resulting list (`2001.0` in the example, above) is the label. The remaining values (`0.884` , `0.610` , etc., in the example) are the features.

After splitting each row, map it to a `LabeledPoint`. You'll have to step down to an RDD (using `.rdd`) or use a DataFrame user-defined function to convert to the `LabeledPoint` object. (See **Hint**, below.) If you step down to an RDD, you'll have to use `toDF()` to convert back to a DataFrame.

Use this new `parse_points` function to parse `raw_data_df`. Then print out the features and label for the first training point, using the `features` and `label` attributes. Finally, calculate the number of features for this dataset.

Hint: Running Arbitrary Lambdas on a DataFrame

To solve this problem, you need a way to run your `parse_points` function on a DataFrame. There are two ways to do this, which we will illustrate with an extremely simple example.

Suppose you have a DataFrame consisting of a first name and a last name, and you want to add a unique [SHA-256](#) hash to each row.

```
df = spark.createDataFrame([('John', "Smith"), ("Ravi",  
"Singh"), ("Julia", "Jones")], ("first_name",  
"last_name"))
```

Here's a simple function to calculate such a hash, using Python's built-in `hashlib` library:

```
def make_hash(first_name, last_name):  
    import hashlib  
    m = hashlib.sha256()  
    # Join the first name and last name by a blank and  
    hash the resulting  
    # string.  
    full_name = ' '.join((first_name, last_name))  
    m.update(full_name)  
    return m.hexdigest()
```

Okay, that's great. But, how do we use it on our DataFrame? We can use a UDF:

```
from pyspark.sql.functions import udf  
u_make_hash = udf(make_hash)  
df2 = df.select(df['*'], u_make_hash(df['first_name'],  
df['last_name']))  
# could run df2.show() here to prove it works
```

Or we can step down to an RDD, use a lambda to call `make_hash` and have the lambda return a `Row` object, which Spark can use to "infer" a new DataFrame.

```
from pyspark.sql import Row
def make_hash_from_row(row):
    hash = make_hash(row[0], row[1])
    return Row(first_name=row[0], last_name=row[1],
               hash=hash)

df2 = (df.rdd
       .map(lambda row: make_hash_from_row(row))
       .toDF())
```

These methods are roughly equivalent. You'll need to do something similar to convert your `raw_data_df` DataFrame into a new DataFrame of `LabeledPoint` objects.

```
In [9]: from pyspark.mllib.regression import LabeledPoint
import numpy as np

# Here is a sample raw data point:
# '2001.0,0.884,0.610,0.600,0.474,0.247,0.357,0.344,0.33,0.600,0.425,0.60
# In this raw data point, 2001.0 is the label, and the remaining values a
```

```
In [14]: # TODO: Replace <FILL IN> with appropriate code
from pyspark.sql import functions as sql_functions
from pyspark.ml.linalg import Vectors

def parse_points(df):
    """Converts a DataFrame of comma separated unicode strings into a DataFrame where each row is a LabeledPoint.
    Args:
        df: DataFrame where each row is a comma separated unicode string.
             is the label and the remaining elements are the features.
    Returns:
        DataFrame: Each row is converted into a `LabeledPoint`, which contains the label and the features. To convert an RDD to a DataFrame, simply call toDF()
    """
    def parse_line(row):
        raw_string = row[0]
        values = [float(x) for x in raw_string.split(',')]

        return LabeledPoint(values[0], values[1:])

    return df.rdd.map(parse_line).toDF()

parsed_points_df = parse_points(raw_data_df)
first_point = parsed_points_df.first()
first_point_features = first_point.features
first_point_label = first_point.label

print(first_point_features, first_point_label)

d = len(first_point_features)

print(d)
```

```
[0.884123733793, 0.610454259079, 0.600498416968, 0.474669212493, 0.24723268094
7, 0.357306088914, 0.344136412234, 0.339641227335, 0.600858840135, 0.4257046890
24, 0.60491501652, 0.419193351817] 2001.0
12
Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in
manager
BrokenPipeError: [Errno 32] Broken pipe
```

In [15]:

```
# TEST Using LabeledPoint (1b)
testmti850.Test.assertTrue(isinstance(first_point_label, float), 'label m
expectedX0 = [0.8841, 0.6105, 0.6005, 0.4747, 0.2472, 0.3573, 0.3441, 0.3396, 0.6
testmti850.Test.assertTrue(np.allclose(expectedX0, first_point_features,
                                         'incorrect features for firstPointFeatures')
testmti850.Test.assertTrue(np.allclose(2001.0, first_point_label), 'incor
testmti850.Test.assertTrue(d == 12, 'incorrect number of features')
```

1 test passed.
1 test passed.
1 test passed.
1 test passed.

Visualization 1: Features

First we will load and setup the visualization library. Then we will look at the raw features for 50 data points by generating a heatmap that visualizes each feature on a grey-scale and shows the variation of each feature across the 50 sample data points. The features are all between 0 and 1, with values closer to 1 represented via darker shades of grey.

In [16]:

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from utilmti850 import prepareSubplot

# takeSample(withReplacement, num, [seed]) randomly selects num elements
# optional seed parameter that one can set for reproducible results

data_values = (parsed_points_df
               .rdd
               .map(lambda lp: lp.features.toArray())
               .takeSample(False, 50, 47))

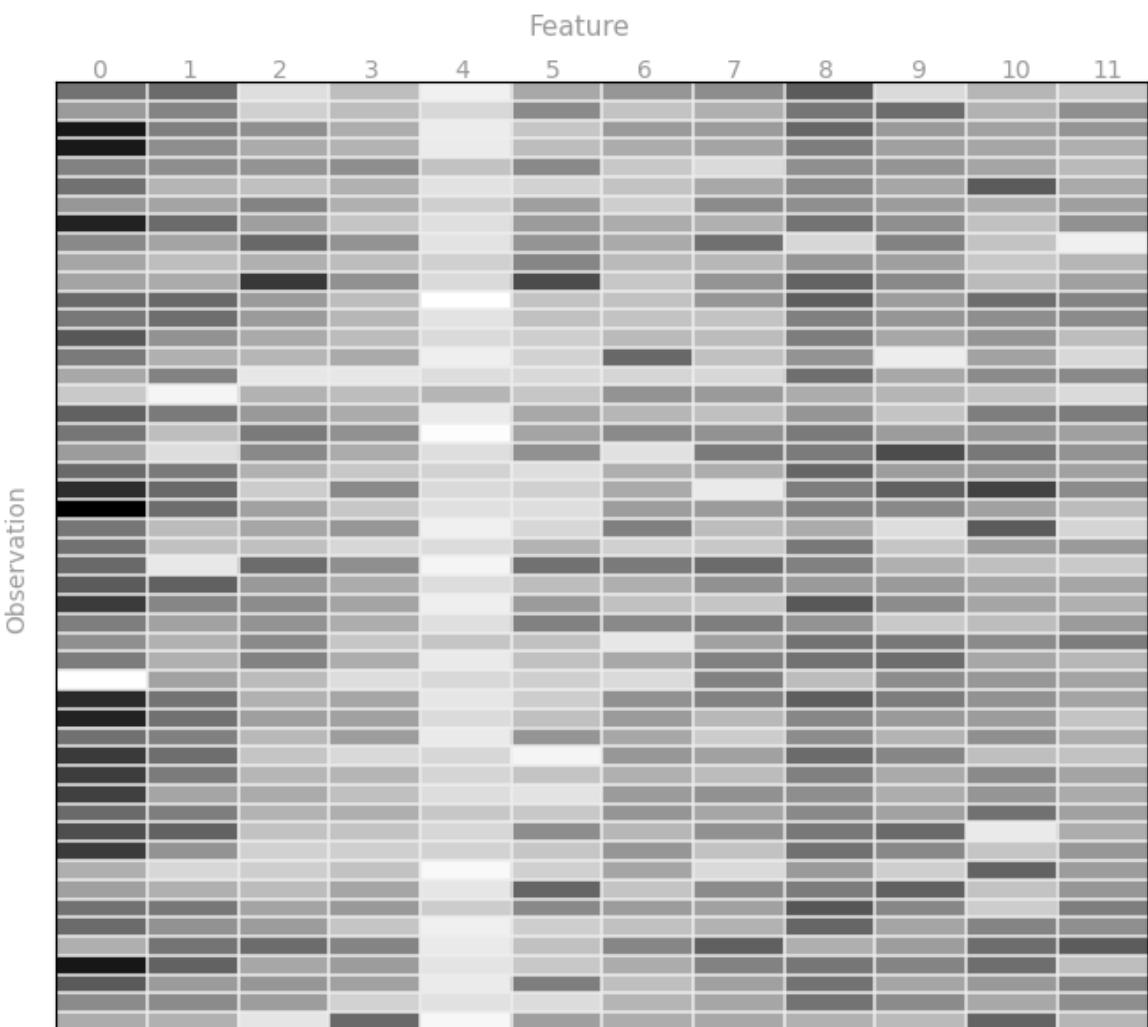
# You can uncomment the line below to see randomly selected features. Th
# selected each time you run the cell because there is no set seed. Note
# this cell with the line commented out when answering the lab quiz quest
# data_values = (parsedPointsDF
#                 .rdd
#                 .map(lambda lp: lp.features.toArray())
#                 .takeSample(False, 50))

# generate layout and plot
fig, ax = prepareSubplot(
    np.arange(.5, 11, 1), np.arange(.5, 49, 1), figsize=(8,7),
    hideLabels=True, gridColor='#eeeeee', gridWidth=1.1
)
image = ax.imshow(data_values, interpolation='nearest', aspect='auto', cma
```

```

for x, y, s in zip(np.arange(-.125, 12, 1), np.repeat(-.75, 12), [str(x)
    ax.text(x, y, s, color='#999999', size='10')
ax.text(4.7, -3, 'Feature', color='red', size='11'), ax.set_ylabel('0
plt.show()

```



(1c) Find the range

Now let's examine the labels to find the range of song years. To do this, find the smallest and largest labels in the `parsed_points_df`.

We will use the `min` and `max` functions that are native to the DataFrames, and thus can be optimized using Spark's Catalyst Optimizer and Project Tungsten (don't worry about the technical details). This code will run faster than simply using the native `min` and `max` functions in Python. Use `selectExpr` to retrieve the `min` and `max` label values.

```
In [26]: # TODO: Replace <FILL IN> with appropriate code
content_stats = parsed_points_df.selectExpr("min(label) as y_min", "max(l
min_year = content_stats['y_min']
max_year = content_stats['y_max']
print(min_year, max_year)
```

1922.0 2011.0

```
In [27]: # TEST Find the range (1c)
testmti850.Test.assertEquals(len(parsed_points_df.first().features), 12,
                             'unexpected number of features in sample point')
sum_feat_two = parsed_points_df.rdd.map(lambda lp: lp.features[2]).sum()
testmti850.Test.assertTrue(np.allclose(sum_feat_two, 3158.96224351), 'par
year_range = max_year - min_year
testmti850.Test.assertTrue(year_range == 89, 'incorrect range for minYear

1 test passed.

Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in
manager
BrokenPipeError: [Errno 32] Broken pipe
1 test passed.
1 test passed.
```

(1d) Shift labels

As we just saw, the labels are years in the 1900s and 2000s. In learning problems, it is often natural to shift labels such that they start from zero. Starting with `parsed_points_df`, create a new DataFrame in which the labels are shifted such that smallest label equals zero (hint: use `select`). After, use `withColumnRenamed` to rename the appropriate columns to `features` and `label`.

```
In [33]: # TODO: Replace <FILL IN> with appropriate code
parsed_data_df = parsed_points_df.select((parsed_points_df.label-min_year

# View the first point
print('\n{0}'.format(parsed_data_df.first()))

Row(label=79.0, features=DenseVector([0.8841, 0.6105, 0.6005, 0.4747, 0.24
72, 0.3573, 0.3441, 0.3396, 0.6009, 0.4257, 0.6049, 0.4192])))

Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in
manager
BrokenPipeError: [Errno 32] Broken pipe
```

```
In [34]: # TEST Shift labels (1d)
old_sample_features = parsed_points_df.first().features
new_sample_features = parsed_data_df.first().features
testmti850.Test.assertTrue(np.allclose(old_sample_features, new_sample_fe
                           'new features do not match old features')
sum_feat_two = parsed_data_df.rdd.map(lambda lp: lp.features[2]).sum()
testmti850.Test.assertTrue(np.allclose(sum_feat_two, 3158.96224351), 'par
min_year_new = parsed_data_df.groupBy().min('label').first()[0]
max_year_new = parsed_data_df.groupBy().max('label').first()[0]
testmti850.Test.assertTrue(min_year_new == 0, 'incorrect min year in shif
testmti850.Test.assertTrue(max_year_new == 89, 'incorrect max year in shi
```

```
Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in
manager
BrokenPipeError: [Errno 32] Broken pipe
Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in
manager
BrokenPipeError: [Errno 32] Broken pipe
1 test passed.

1 test passed.
1 test passed.
1 test passed.
```

Visualization 2: Shifting labels

We will look at the labels before and after shifting them. Both scatter plots below visualize tuples storing:

- a label value and
- the number of training points with this label.

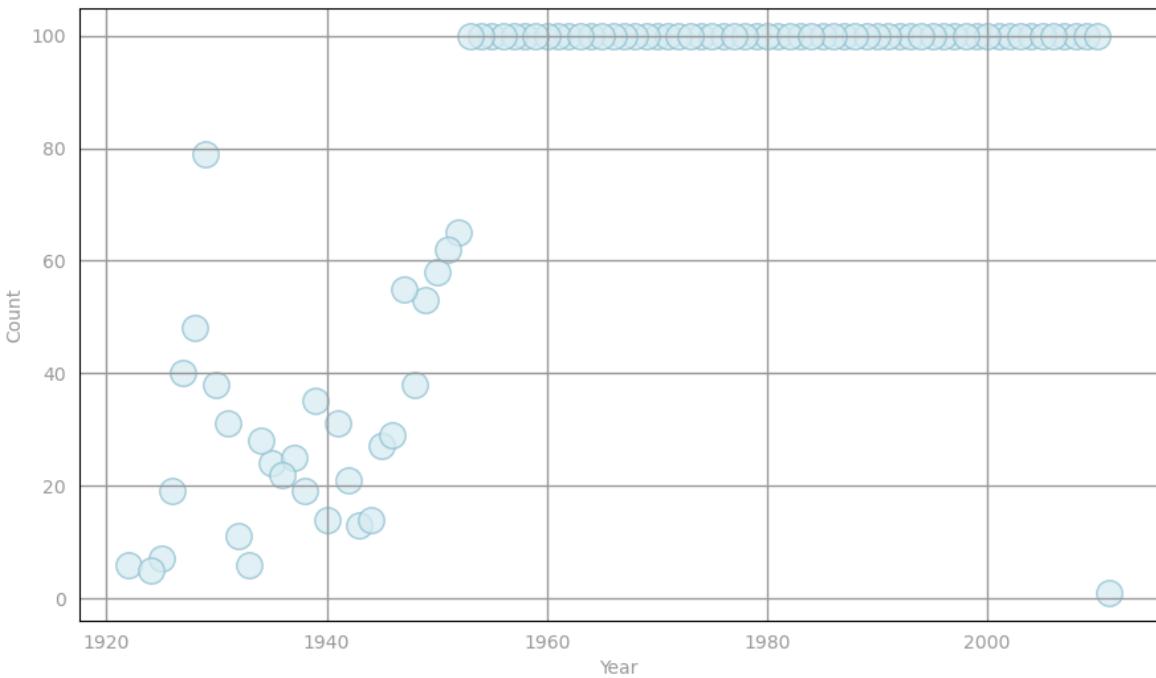
The first scatter plot uses the initial labels, while the second one uses the shifted labels. Note that the two plots look the same except for the labels on the x-axis.

In [35]:

```
# from operator import add

# get data for plot
old_data = (parsed_points_df
            .rdd
            .map(lambda lp: (lp.label, 1))
            .reduceByKey(lambda x, y: x + y) # alternatively, one can us
            .collect())                      # instead of the lambda fun
x, y = zip(*old_data)

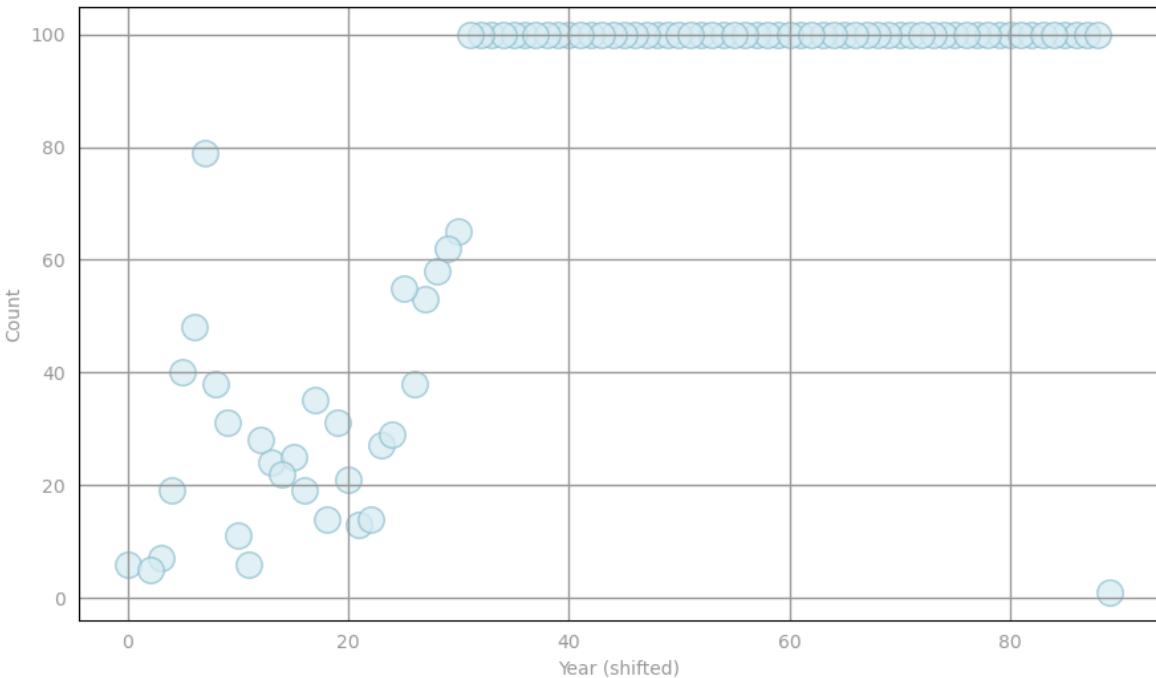
# generate layout and plot data
fig, ax = prepareSubplot(np.arange(1920, 2050, 20), np.arange(0, 150, 20))
ax.scatter(x, y, s=14**2, c='#d6ebf2', edgecolors='#8cbfd0', alpha=0.75)
ax.set_xlabel('Year')
ax.set_ylabel('Count')
plt.show()
```



In [36]:

```
# get data for plot
new_data = (parsed_data_df
    .rdd
    .map(lambda lp: (lp.label, 1))
    .reduceByKey(lambda x, y: x + y)
    .collect())
x, y = zip(*new_data)

# generate layout and plot data
fig, ax = prepareSubplot(np.arange(0, 120, 20), np.arange(0, 120, 20))
ax.scatter(x, y, s=14**2, c='#d6ebf2', edgecolors='#8cbfd0', alpha=0.75)
ax.set_xlabel('Year (shifted)')
ax.set_ylabel('Count')
plt.show()
```



(1e) Training, validation, and test sets

We're almost done parsing our dataset, and our final task involves splitting the dataset into training, validation and test sets. Use the `randomSplit` method with the specified weights and seed to create DataFrames storing each of these datasets. Next, cache each of these DataFrames, as we will be accessing them multiple times in the remainder of this lab. Finally, compute the size of each dataset and verify that the sum of their sizes equals the value computed in Part (1a).

```
In [37]: # TODO: Replace <FILL IN> with appropriate code
weights = [.8, .1, .1]
seed = 42

parsed_train_data_df, parsed_val_data_df, parsed_test_data_df = parsed_da
parsed_train_data_df.cache
parsed_val_data_df.cache
parsed_test_data_df.cache
n_train = parsed_train_data_df.count()
n_val = parsed_val_data_df.count()
n_test = parsed_test_data_df.count()
print(n_train, n_val, n_test, n_train + n_val + n_test)
print(parsed_data_df.count())
```

5437 612 675 6724
6724

```
In [38]: # TEST Training, validation, and test sets (1e)
testmti850.Test.assertEquals(len(parsed_train_data_df.first().features),
                             'parsed_train_data_df has wrong number of features')
sum_feat_two = (parsed_train_data_df
                 .rdd
                 .map(lambda lp: lp.features[2])
                 .sum())
sum_feat_three = (parsed_val_data_df
                   .rdd
                   .map(lambda lp: lp.features[3])
                   .reduce(lambda x, y: x + y))
sum_feat_four = (parsed_test_data_df
                  .rdd
                  .map(lambda lp: lp.features[4])
                  .reduce(lambda x, y: x + y))
testmti850.Test.assertTrue(np.allclose([sum_feat_two, sum_feat_three, sum
                                         2526.87757656, 297.340394298, 184.235876654],
                                         'parsed Train, Val, Test data has unexpected values'))
testmti850.Test.assertTrue(n_train + n_val + n_test == 6724, 'unexpected')
testmti850.Test.assertEquals(n_train, 5437, 'unexpected value for nTrain')
testmti850.Test.assertEquals(n_val, 612, 'unexpected value for nVal')
testmti850.Test.assertEquals(n_test, 675, 'unexpected value for nTest')
```

```
1 test passed.  
1 test passed.
```

Part 2: Create and evaluate a baseline model

(2a) Average label

A very simple yet natural baseline model is one where we always make the same prediction independent of the given data point, using the average label in the training set as the constant prediction value. Compute this value, which is the average (shifted) song year for the training set. Use `selectExpr` and `first()` from the [DataFrame API](#).

```
In [44]: # TODO: Replace <FILL IN> with appropriate code  
  
average_train_year = parsed_train_data_df.selectExpr('avg(label)').first()  
  
print(average_train_year)  
  
53.54919992643001
```

```
In [45]: # TEST Average label (2a)  
testmti850.Test.assertTrue(  
    np.allclose(average_train_year, 53.54919992643001), 'incorrect value'  
)  
  
1 test passed.
```

(2b) Root mean squared error

We naturally would like to see how well this naive baseline performs. We will use root mean squared error ([RMSE](#)) for evaluation purposes. Using [Regression Evaluator](#), compute the RMSE given a dataset of $(prediction, label)$ tuples.

```
In [41]: # TODO: Replace <FILL IN> with appropriate code  
from pyspark.ml.evaluation import RegressionEvaluator  
  
preds_and_labels = [(1., 3.), (2., 1.), (2., 2.)]  
  
preds_and_labels_df = spark.createDataFrame(preds_and_labels, ["prediction", "label"])  
  
evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="label")  
  
def calc_RMSE(dataset):  
    """Calculates the root mean squared error for an dataset of (prediction, label) tuples.  
  
    Args:  
        dataset (DataFrame of (float, float)): A `DataFrame` consisting of two columns.  
        The first column is labeled "prediction" and the second column is labeled "label".  
  
    Returns:  
        float: The square root of the mean of the squared errors.  
    """  
  
    rmse = evaluator.evaluate(dataset)  
  
    return rmse
```

```

    """
    return evaluator.evaluate(dataset)

example_rmse = calc_RMSE(preds_and_labels_df)

print(example_rmse)
# RMSE = sqrt[((1-3)^2 + (2-1)^2 + (2-2)^2) / 3] = 1.291
1.2909944487358058

```

In [42]: `testmti850.Test.assertTrue(np.allclose(example_rmse, 1.2909944487358058),
1 test passed.`

(2c) Training, validation and test RMSE

Now let's calculate the training, validation and test RMSE of our baseline model. To do this, first create DataFrames of (*prediction, label*) tuples for each dataset, and then call `calc_RMSE()`. Note that each RMSE can be interpreted as the average prediction error for the given dataset (in terms of number of years). You can use `createDataFrame` to make a DataFrame with the column names of "prediction" and "label" from an RDD.

In [46]: `# TODO: Replace <FILL IN> with appropriate code`

```

preds_and_labels_train = parsed_train_data_df.rdd.map(lambda row: (average_
preds_and_labels_train_df = spark.createDataFrame(preds_and_labels_train,
rmse_train_base = calc_RMSE(preds_and_labels_train_df)

preds_and_labels_val = parsed_val_data_df.rdd.map(lambda row: (average_
preds_and_labels_val_df = spark.createDataFrame(preds_and_labels_val, ["p
rmse_val_base = calc_RMSE(preds_and_labels_val_df)

preds_and_labels_test = parsed_test_data_df.rdd.map(lambda row: (average_
preds_and_labels_test_df = spark.createDataFrame(preds_and_labels_test, [
rmse_test_base = calc_RMSE(preds_and_labels_test_df)

print('Baseline Train RMSE = {:.3f}'.format(rmse_train_base))
print('Baseline Validation RMSE = {:.3f}'.format(rmse_val_base))
print('Baseline Test RMSE = {:.3f}'.format(rmse_test_base))

Baseline Train RMSE = 21.598
Baseline Validation RMSE = 20.759
Baseline Test RMSE = 20.547

```

In [48]: `# TEST Training, validation and test RMSE (2c)`

```

testmti850.Test.assertTrue(
    np.allclose(
        [rmse_train_base, rmse_val_base, rmse_test_base],
        [21.598132, 20.759189, 20.546997]
    ),

```

```
    'incorrect RMSE values'
)
```

1 test passed.

Visualization 3: Predicted vs. actual

We will visualize predictions on the validation dataset. The scatter plots below visualize tuples storing i) the predicted value and ii) true label. The first scatter plot represents the ideal situation where the predicted value exactly equals the true label, while the second plot uses the baseline predictor (i.e., `average_train_year`) for all predicted values. Further note that the points in the scatter plots are color-coded, ranging from light yellow when the true and predicted values are equal to bright red when they drastically differ.

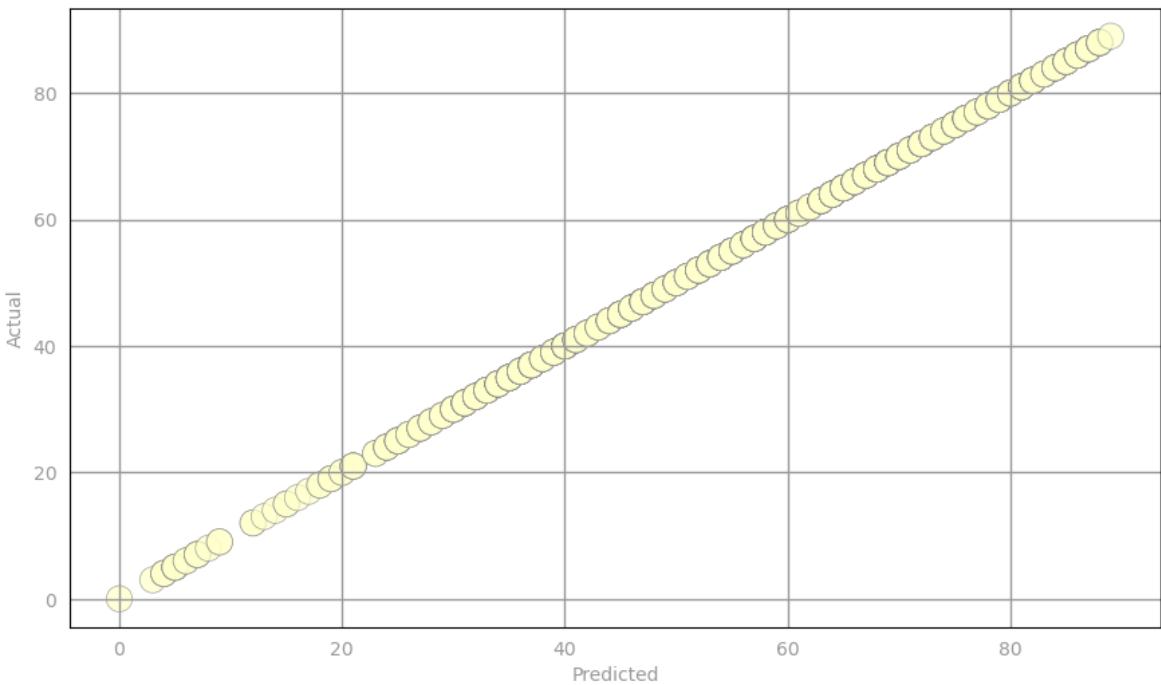
```
In [49]: from matplotlib.colors import ListedColormap, Normalize
from matplotlib.cm import get_cmap
cmap = get_cmap('YlOrRd')
norm = Normalize()

def squared_error(label, prediction):
    """Calculates the squared error for a single prediction."""
    return float((label - prediction)**2)

actual = np.asarray(parsed_val_data_df
                    .select('label')
                    .collect())
error = np.asarray(parsed_val_data_df
                    .rdd
                    .map(lambda lp: (lp.label, lp.label))
                    .map(lambda lp: squared_error(lp[0], lp[1]))
                    .collect())
clrs = cmap(np.asarray(norm(error))[:, :3]

fig, ax = prepareSubplot(np.arange(0, 100, 20), np.arange(0, 100, 20))
ax.scatter(actual, actual, s=14**2, c=clrs, edgecolors="#888888", alpha=0
ax.set_xlabel('Predicted')
ax.set_ylabel('Actual')
plt.show()
```

/tmp/ipykernel_22549/2990054735.py:3: MatplotlibDeprecationWarning: The `get_cmap` function was deprecated in Matplotlib 3.7 and will be removed in 3.11. Use `matplotlib.colormaps[name]` or `matplotlib.colormaps.get_cmap()` or `pyplot.get_cmap()` instead.
 cmap = get_cmap('YlOrRd')



```
In [50]: def squared_error(label, prediction):
    """Calculates the squared error for a single prediction."""
    return float((label - prediction) ** 2)

predictions = np.asarray(parsed_val_data_df
                        .rdd
                        .map(lambda lp: average_train_year)
                        .collect())
error = np.asarray(parsed_val_data_df
                    .rdd
                    .map(lambda lp: (lp.label, average_train_year['avg(label')])
                    .map(lambda lp: squared_error(lp[0], lp[1]))
                    .collect())

norm = Normalize()
clrs = cmap(np.asarray(norm(error))[:, : 3])

fig, ax = prepareSubplot(np.arange(53.0, 55.0, 0.5), np.arange(0, 100, 20))
ax.set_xlim(53, 55)
ax.scatter(predictions, actual, s=14**2, c=clrs, edgecolors="#888888", alpha=0.8)
ax.set_xlabel('Predicted')
ax.set_ylabel('Actual')
plt.show()
```

```
25/11/28 01:59:06 ERROR Executor: Exception in task 0.0 in stage 80.0 (TID  
67)  
org.apache.spark.api.python.PythonException: Traceback (most recent call l  
ast):  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/worker.py", line 2044, i  
n main  
    process()  
    ~~~~~^~  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/worker.py", line 2036, i  
n process  
    serializer.dump_stream(out_iter, outfile)  
    ~~~~~^~  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/serializers.py", line 27  
3, in dump_stream  
    vs = list(itertools.islice(iterator, batch))  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/util.py", line 131, in w  
rapper  
    return f(*args, **kwargs)  
File "/tmp/ipykernel_22549/2256473714.py", line 11, in <lambda>  
TypeError: 'float' object is not subscriptable  
  
          at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.han  
dlePythonException(PythonRunner.scala:581)  
          at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRun  
ner.scala:940)  
          at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRun  
ner.scala:925)  
          at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.has  
Next(PythonRunner.scala:532)  
          at org.apache.spark.InterruptibleIterator.hasNext(InterruptibleIte  
rator.scala:37)  
          at scala.collection.mutable.Growable.addAll(Growable.scala:61)  
          at scala.collection.mutable.Growable.addAll$(Growable.scala:57)  
          at scala.collection.mutable.ArrayBuilder.addAll(ArrayBuilder.scala  
a:75)  
          at scala.collection.IterableOnceOps.toArray(IterableOnce.scala:150  
5)  
          at scala.collection.IterableOnceOps.toArray$(IterableOnce.scala:14  
98)  
          at org.apache.spark.InterruptibleIterator.toArray(InterruptibleIte  
rator.scala:28)  
          at org.apache.spark.rdd.RDD.$anonfun$collect$2(RDD.scala:1057)  
          at org.apache.spark.SparkContext.$anonfun$runJob$5(SparkContext.sc  
ala:2524)  
          at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala  
a:93)  
          at org.apache.spark.TaskContext.runTaskWithListeners(TaskContext.s  
cala:171)  
          at org.apache.spark.scheduler.Task.run(Task.scala:147)  
          at org.apache.spark.executor.Executor$TaskRunner.$anonfun$run$5(Ex  
ecutor.scala:647)  
          at org.apache.spark.util.SparkErrorUtils.tryWithSafeFinally(SparkE  
rrorUtils.scala:80)  
          at org.apache.spark.util.SparkErrorUtils.tryWithSafeFinally$(SparkE  
rrorUtils.scala:77)  
          at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:99)  
          at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala  
a:650)  
          at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(Thread  
PoolExecutor.java:1144)
```



```
        at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1144)
        at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:642)
        at java.base/java.lang.Thread.run(Thread.java:1583)
```

25/11/28 01:59:06 ERROR TaskSetManager: Task 0 in stage 80.0 failed 1 time
s; aborting job

```
- Py4JJavaError                                     Traceback (most recent call last)
t)
Cell In[50], line 13
    3     return float((label - prediction) ** 2)
    5 predictions = np.asarray(parsed_val_data_df
    6             .rdd
    7             .map(lambda lp: average_train_year)
    8             .collect())
    9 error = np.asarray(parsed_val_data_df
   10            .rdd
   11            .map(lambda lp: (lp.label, average_train_year
['avg(label')]))
   12            .map(lambda lp: squared_error(lp[0], lp[1])))
--> 13            .collect())
   15 norm = Normalize()
   16 clrs = cmap(np.asarray(norm(error))[:, : 3]

File /opt/spark/python/pyspark/core/rdd.py:1700, in RDD.collect(self)
 1698 with SCCallSiteSync(self.context):
 1699     assert self.ctx._jvm is not None
-> 1700     sock_info = self.ctx._jvm.PythonRDD.collectAndServe(self._jrdd.rdd())
 1701 return list(_load_from_socket(sock_info, self._jrdd_deserializer))

File /opt/spark/python/lib/py4j-0.10.9.9-src.zip/py4j/java_gateway.py:136
2, in JavaMember.__call__(self, *args)
 1356 command = proto.CALL_COMMAND_NAME +\
 1357     self.command_header +\
 1358     args_command +\
 1359     proto.END_COMMAND_PART
 1361 answer = self.gateway_client.send_command(command)
-> 1362 return_value = get_return_value(
 1363     answer, self.gateway_client, self.target_id, self.name)
 1365 for temp_arg in temp_args:
 1366     if hasattr(temp_arg, "_detach"):

File /opt/spark/python/pyspark/errors/exceptions/captured.py:282, in capture_sql_exception.<locals>.deco(*a, **kw)
 279 from py4j.protocol import Py4JJavaError
 281 try:
--> 282     return f(*a, **kw)
 283 except Py4JJavaError as e:
 284     converted = convert_exception(e.java_exception)

File /opt/spark/python/lib/py4j-0.10.9.9-src.zip/py4j/protocol.py:327, in get_return_value(answer, gateway_client, target_id, name)
 325 value = OUTPUT_CONVERTER[type](answer[2:], gateway_client)
 326 if answer[1] == REFERENCE_TYPE:
--> 327     raise Py4JJavaError(
 328         "An error occurred while calling {0}{1}{2}.\\n".
 329         format(target_id, ".", name), value)
 330 else:
 331     raise Py4JError(
 332         "An error occurred while calling {0}{1}{2}. Trace:\\n{3}\\n".
 333         format(target_id, ".", name, value))

Py4JJavaError: An error occurred while calling z:org.apache.spark.api.pyth
```

```
on.PythonRDD.collectAndServe.  
: org.apache.spark.SparkException: Job aborted due to stage failure: Task  
0 in stage 80.0 failed 1 times, most recent failure: Lost task 0.0 in stag  
e 80.0 (TID 67) (10.0.2.15 executor driver): org.apache.spark.api.python.P  
ythonException: Traceback (most recent call last):  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/worker.py", line 2044, i  
n main  
    process()  
    ~~~~~^~  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/worker.py", line 2036, i  
n process  
    serializer.dump_stream(out_iter, outfile)  
    ~~~~~^~~~~~  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/serializers.py", line 27  
3, in dump_stream  
    vs = list(itertools.islice(iterator, batch))  
  File "/opt/spark/python/lib/pyspark.zip/pyspark/util.py", line 131, in w  
rapper  
    return f(*args, **kwargs)  
File "/tmp/ipykernel_22549/2256473714.py", line 11, in <lambda>  
TypeError: 'float' object is not subscriptable  
  
      at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.han  
dlePythonException(PythonRunner.scala:581)  
      at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRun  
ner.scala:940)  
      at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRun  
ner.scala:925)  
      at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.has  
Next(PythonRunner.scala:532)  
      at org.apache.spark.InterruptibleIterator.hasNext(InterruptibleIte  
rator.scala:37)  
      at scala.collection.mutable.Growable.addAll(Growable.scala:61)  
      at scala.collection.mutable.Growable.addAll$(Growable.scala:57)  
      at scala.collection.mutable.ArrayBuilder.addAll(ArrayBuilder.scala  
a:75)  
      at scala.collection.IterableOnceOps.toArray(IterableOnce.scala:150  
5)  
      at scala.collection.IterableOnceOps.toArray$(IterableOnce.scala:14  
98)  
      at org.apache.spark.InterruptibleIterator.toArray(InterruptibleIte  
rator.scala:28)  
      at org.apache.spark.rdd.RDD.$anonfun$collect$2(RDD.scala:1057)  
      at org.apache.spark.SparkContext.$anonfun$runJob$5(SparkContext.sc  
ala:2524)  
      at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala  
a:93)  
      at org.apache.spark.TaskContext.runTaskWithListeners(TaskContext.s  
cala:171)  
      at org.apache.spark.scheduler.Task.run(Task.scala:147)  
      at org.apache.spark.executor.Executor$TaskRunner.$anonfun$run$5(Ex  
ecutor.scala:647)  
      at org.apache.spark.util.SparkErrorUtils.tryWithSafeFinally(SparkE  
rrorUtils.scala:80)  
      at org.apache.spark.util.SparkErrorUtils.tryWithSafeFinally$(Spark  
ErrorUtils.scala:77)  
      at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:99)  
      at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala  
a:650)  
      at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(Thr
```

```
eadPoolExecutor.java:1144)
    at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:642)
    at java.base/java.lang.Thread.run(Thread.java:1583)

Driver stacktrace:
    at org.apache.spark.scheduler.DAGScheduler.$anonfun$abortStage$3(DAGScheduler.scala:2935)
    at scala.Option.getOrElse(Option.scala:201)
    at org.apache.spark.scheduler.DAGScheduler.$anonfun$abortStage$2(DAGScheduler.scala:2935)
    at org.apache.spark.scheduler.DAGScheduler.$anonfun$abortStage$2$adapted(DAGScheduler.scala:2927)
    at scala.collection.immutable.List.foreach(List.scala:334)
    at org.apache.spark.scheduler.DAGScheduler.abortStage(DAGScheduler.scala:2927)
    at org.apache.spark.scheduler.DAGScheduler.$anonfun$handleTaskSetFailed$1(DAGScheduler.scala:1295)
    at org.apache.spark.scheduler.DAGScheduler.$anonfun$handleTaskSetFailed$1$adapted(DAGScheduler.scala:1295)
    at scala.Option.foreach(Option.scala:437)
    at org.apache.spark.scheduler.DAGScheduler.handleTaskSetFailed(DAGScheduler.scala:1295)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.doOnReceive(DAGScheduler.scala:3207)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.onReceive(DAGScheduler.scala:3141)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.onReceive(DAGScheduler.scala:3130)
    at org.apache.spark.util.EventLoop$$anon$1.run(EventLoop.scala:50)
    at org.apache.spark.scheduler.DAGScheduler.runJob(DAGScheduler.scala:1009)
    at org.apache.spark.SparkContext.runJob(SparkContext.scala:2484)
    at org.apache.spark.SparkContext.runJob(SparkContext.scala:2505)
    at org.apache.spark.SparkContext.runJob(SparkContext.scala:2524)
    at org.apache.spark.SparkContext.runJob(SparkContext.scala:2549)
    at org.apache.spark.rdd.RDD.$anonfun$collect$1(RDD.scala:1057)
    at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:151)
    at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.scala:112)
    at org.apache.spark.rdd.RDD.withScope(RDD.scala:417)
    at org.apache.spark.rdd.RDD.collect(RDD.scala:1056)
    at org.apache.spark.api.python.PythonRDD$.collectAndServe(PythonRDD.scala:203)
    at org.apache.spark.api.python.PythonRDD.collectAndServe(PythonRDD.scala)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:75)
    at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:52)
    at java.base/java.lang.reflect.Method.invoke(Method.java:580)
    at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
    at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:37)
    at py4j.Gateway.invoke(Gateway.java:282)
    at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
```



```

        at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:650)
        at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1144)
        at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:642)
        ... 1 more

```

Part 3: Train (via gradient descent) and evaluate a linear regression model

(3a) Gradient summand

Now let's see if we can do better via linear regression, training a model via gradient descent (we'll omit the intercept for now). Recall that the gradient descent update for linear regression is:

$$\mathbf{w}_{i+1} = \mathbf{w}_i - \alpha_i \sum_j (\mathbf{w}_i^\top \mathbf{x}_j - y_j) \mathbf{x}_j.$$

where i is the iteration number of the gradient descent algorithm, and j identifies the observation.

First, implement a function that computes the summand for this update, i.e., the summand equals $(\mathbf{w}^\top \mathbf{x} - y)\mathbf{x}$, and test out this function on two examples. Use the `DenseVector` `dot` method.

```
In [51]: from pyspark.mllib.linalg import DenseVector

# OBS: You are advised to use pyspark.ml.linalg.DenseVector and adapt the

In [55]: # TODO: Replace <FILL IN> with appropriate code
def gradient_summand(weights, lp):
    """Calculates the gradient summand for a given weight and `LabeledPoint`"""

    Note:
        `DenseVector` behaves similarly to a `numpy.ndarray` and they can
        be used interchangeably within this function. For example, they both implement the `dot` method.

    Args:
        weights (DenseVector): An array of model weights (betas).
        lp (LabeledPoint): The `LabeledPoint` for a single observation.

    Returns:
        DenseVector: An array of values the same length as `weights`. The
        summand  $(\mathbf{w}^\top \mathbf{x} - y)\mathbf{x}$ .
    """
    return (weights.dot(lp.features) - lp.label) * (lp.features)

example_w = DenseVector([1, 1, 1])
example_lp = LabeledPoint(2.0, [3, 1, 4])
# gradient_summand = (dot([1 1 1], [3 1 4]) - 2) * [3 1 4] = (8 - 2) * [3 1 4]
summand_one = gradient_summand(example_w, example_lp)
print(summand_one)
```

```

example_w = DenseVector([.24, 1.2, -1.4])
example_lp = LabeledPoint(3.0, [-1.4, 4.2, 2.1])
summand_two = gradient_summand(example_w, example_lp)
print(summand_two)

[18.0, 6.0, 24.0]
[1.7304000000000002, -5.191200000000001, -2.5956000000000006]

```

In [56]: # TEST Gradient summand (3a)

```

testmti850.Test.assertTrue(np.allclose(summand_one, [18., 6., 24.]), 'incorrect')
testmti850.Test.assertTrue(np.allclose(summand_two, [1.7304, -5.1912, -2.5956]), 'incorrect')

1 test passed.
1 test passed.

```

(3b) Use weights to make predictions

Next, implement a `get_labeled_predictions` function that takes in weights and an observation's `LabeledPoint` and returns a `(prediction, label)` tuple. Note that we can predict by computing the dot product between weights and an observation's features.

In [60]: # TODO: Replace <FILL IN> with appropriate code

```

def get_labeled_prediction(weights, observation):
    """Calculates predictions and returns a (prediction, label) tuple.

    Note:
        The labels should remain unchanged as we'll use this information
        error later.

    Args:
        weights (np.ndarray): An array with one weight for each features
        observation (LabeledPoint): A `LabeledPoint` that contain the correct
            features for the data point.

    Returns:
        tuple: A (prediction, label) tuple. Convert the return type of this
        """
    return (float(weights.dot(observation.features)), float(observation.label))

weights = np.array([1.0, 1.5])

prediction_example = spark.sparkContext.parallelize([LabeledPoint(2, np.array([-1.4, 4.2, 2.1])), LabeledPoint(1.5, np.array([.5, .5]))])

preds_and_labels_example = prediction_example.map(lambda lp: get_labeled_prediction(lp))

print(preds_and_labels_example.collect())

```

[(1.75, 2.0), (1.25, 1.5)]

```
In [61]: # TEST Use weights to make predictions (3b)
testmti850.Test.assertTrue(isinstance(preds_and_labels_example.first()[0]
testmti850.Test.assertEquals(preds_and_labels_example.collect(), [(1.75,
                           'incorrect definition for getLabeledPredictions')]
```

```
1 test passed.
1 test passed.
```

(3c) Gradient descent

Next, implement a gradient descent function for linear regression and test out this function on an example.

```
In [65]: # TODO: Replace <FILL IN> with appropriate code

def linreg_gradient_descent(train_data, num_iters):
    """Calculates the weights and error for a linear regression model trained
    on the provided training data.

    Note:
        `DenseVector` behaves similarly to a `numpy.ndarray` and they can
        be used interchangeably within this function. For example, they both implement the `dot` method.

    Args:
        train_data (RDD of LabeledPoint): The labeled data for use in training.
        num_iters (int): The number of iterations of gradient descent to run.

    Returns:
        (np.ndarray, np.ndarray): A tuple of (weights, training errors).
        The first element is a `DenseVector` containing the final weights (one weight per feature) for the model, and the second element is a list of errors (RMSE) for each iteration of the algorithm.
    """
    # The length of the training data
    n = train_data.count()

    # The number of features in the training data
    d = len(train_data.first().features)
    w = np.zeros(d)
    alpha = 1.0

    # We will compute and store the training error after each iteration
    error_train = np.zeros(num_iters)
    for i in range(num_iters):
        # Use get_labeled_prediction from (3b) with trainData to obtain a
        # tuples. Note that the weights all equal 0 for the first iteration
        # have large errors to start.

        preds_and_labels_train = train_data.map(lambda lp: get_labeled_pr
preds_and_labels_train_df = spark.createDataFrame(preds_and_labels_train)
error_train[i] = calc_RMSE(preds_and_labels_train_df)

        # Calculate the `gradient`. Make use of the `gradient_summand` function
        # Note that `gradient` should be a `DenseVector` of length `d`.
        gradient = train_data.map(lambda lp: gradient_summand(w, lp)).reduce(
            lambda a, b: a + b)

        # Update the weights
        w = w - alpha * gradient
```

```

        alpha_i = alpha / (n * np.sqrt(i + 1))

        w -= alpha_i*gradient

    return w, error_train

# create a toy dataset with n = 10, d = 3, and then run 5 iterations of g
# note: the resulting model will not be useful; the goal here is to verify
# linreg_gradient_descent is working properly
example_n = 10
example_d = 3
example_data = (
    spark
    .sparkContext
    .parallelize(parsed_train_data_df.take(example_n))
    .map(lambda lp: LabeledPoint(lp.label, lp.features[0:example_d]))
)
print(example_data.take(2))

example_num_iters = 5
example_weights, example_error_train = linreg_gradient_descent(example_da

print(example_weights)
print(example_error_train)

[LabeledPoint(0.0, [0.581044957651, 0.233509721704, 0.819061699014]), LabeledPoint(0.0, [0.63944087626, 0.65293056486, 0.343546487474])]

[0.59874548 0.16755145 1.19239355]
[1.73205081 1.3096377 1.27962073 1.26123272 1.24682059]
```

In [67]:

```
# TEST Gradient descent (3c)
expected_output = [0.59874548, 0.16755145, 1.19239355]
testmti850.Test.assertTrue(np.allclose(example_weights, expected_output),
expected_error = [1.73205081, 1.3096377, 1.27962073, 1.26123272, 1.24682059]
testmti850.Test.assertTrue(np.allclose(example_error_train, expected_error,
                                         atol=1e-05,
                                         msg='value of exampleErrorTrain is incorrect'))
```

1 test passed.
1 test passed.

(3d) Train the model

Now let's train a linear regression model on all of our training data and evaluate its accuracy on the validation set. Note that the test set will not be used here. If we evaluated the model on the test set, we would bias our final results.

We've already done much of the required work: we computed the number of features in Part (1b); we created the training and validation datasets and computed their sizes in Part (1e); and, we wrote a function to compute RMSE in Part (2b).

In [70]:

```
# TODO: Replace <FILL IN> with appropriate code
num_iters = 50

weights_LR0, error_train_LR0 = linreg_gradient_descent(parsed_train_data_
preds_and_labels = parsed_val_data_df.rdd.map(lambda lp: (get_labeled_pre
```

```

preds_and_labels_df = spark.createDataFrame(preds_and_labels, ["prediction", "label"])
rmse_val_LR0 = calc_RMSE(preds_and_labels_df)
print('Validation RMSE:\n\tBaseline = {0:.3f}\n\tLR0 = {1:.3f}'.format(rmse_val_LR0))

```

Validation RMSE:

Baseline = 20.759
LR0 = 18.371

```
In [71]: # TEST Train the model (3d)
expected_output = [22.63418505, 20.64417482, -0.39044365, 8.37584137, 5.8
                   3.6353787, 9.87147047, 5.80887121, 11.0129383, 3.48020
testmti850.Test.assertTrue(np.allclose(weights_LR0, expected_output), 'in'
1 test passed.
```

Visualization 4: Training error

We will look at the log of the training error as a function of iteration. The first scatter plot visualizes the logarithm of the training error for all 50 iterations. The second plot shows the training error itself, focusing on the final 44 iterations.

```
In [72]: norm = Normalize()

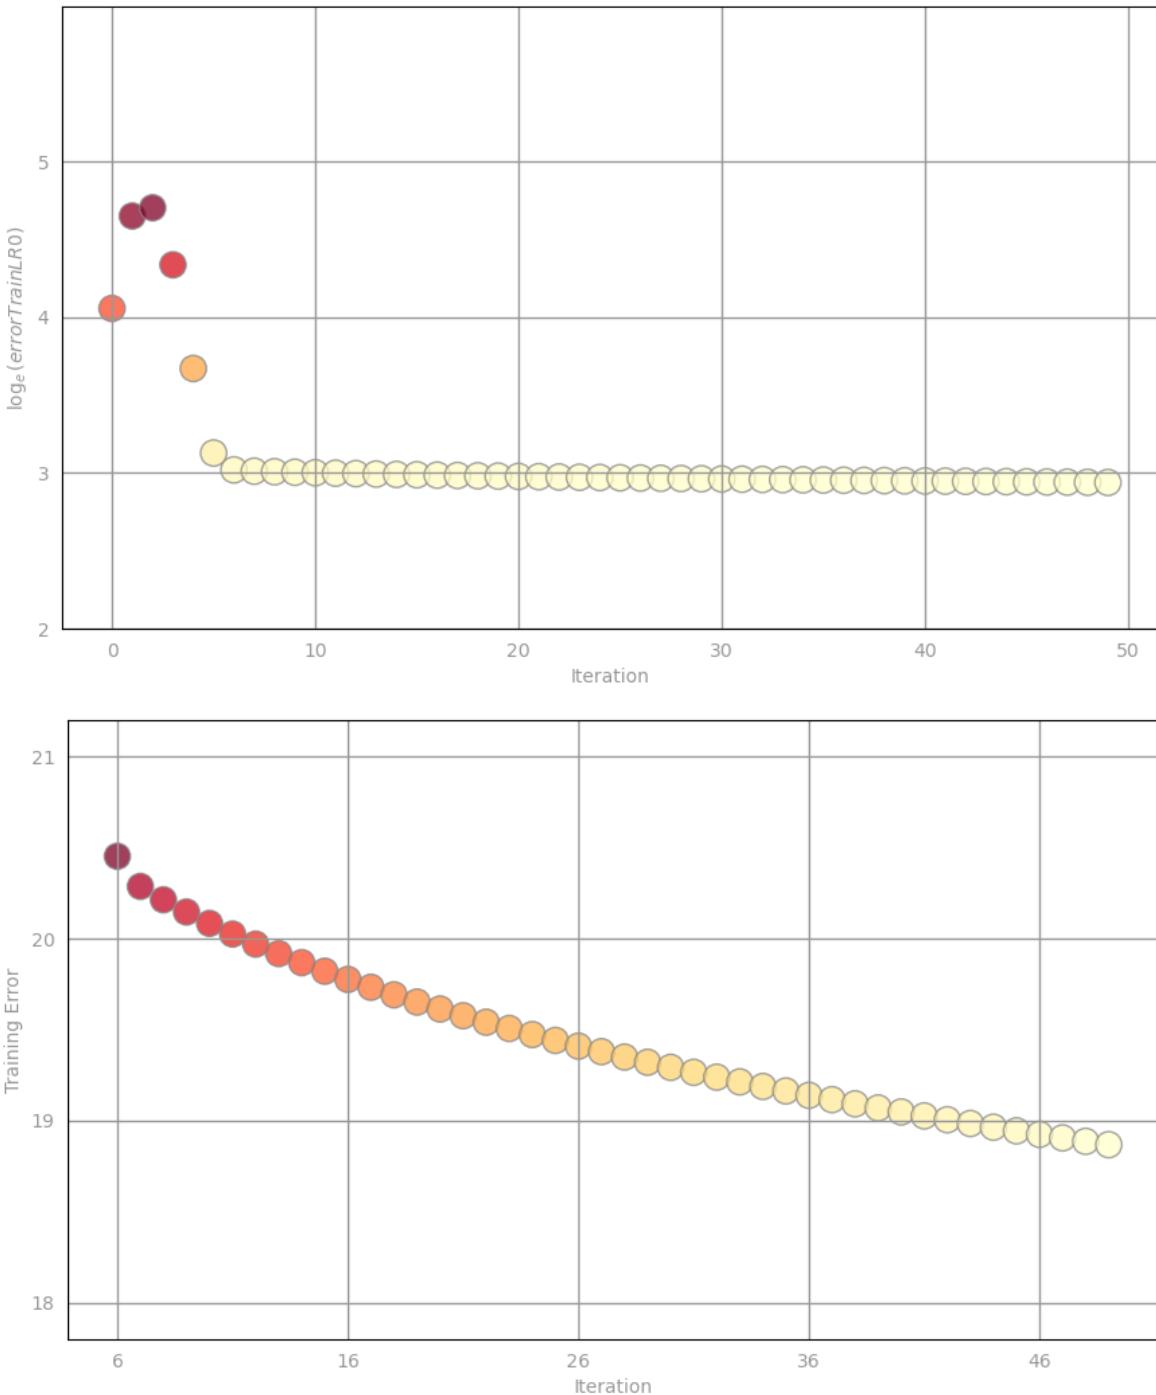
clrs = cmap(np.asarray(norm(np.log(error_train_LR0))))[:,0:3]

fig, ax = prepareSubplot(np.arange(0, 60, 10), np.arange(2, 6, 1))
ax.set_yscale(2, 6)
ax.scatter(range(0, num_iters), np.log(error_train_LR0), s=14**2, c=clrs,
           ax.set_xlabel('Iteration')
           ax.set_ylabel(r'$\log_e(errorTrainLR0)$')
           plt.show((fig))

norm = Normalize()

clrs = cmap(np.asarray(norm(error_train_LR0[6:])))[:,0:3]

fig, ax = prepareSubplot(np.arange(0, 60, 10), np.arange(17, 22, 1))
ax.set_yscale(17.8, 21.2)
ax.scatter(range(0, num_iters - 6), error_train_LR0[6:], s=14**2, c=clrs,
           ax.set_xticklabels(map(str, range(6, 66, 10)))
           ax.set_xlabel('Iteration')
           ax.set_ylabel(r'Training Error')
           plt.show()
```



Part 4: Train using SparkML and perform grid search

(4a) LinearRegression

We're already doing better than the baseline model, but let's see if we can do better by adding an intercept, using regularization, and (based on the previous visualization) training for more iterations. SparkML's [LinearRegression](#) essentially implements the same algorithm that we implemented in Part (3b), albeit more efficiently and with various additional functionality, such as including an intercept in the model and allowing L1, L2, or [elastic net regularization](#). Elastic net regularization is a linear combination of L1 and L2 regularization. For $\alpha = 0$, the penalty is an L2 penalty. For $\alpha = 1$, it is an L1

penalty.

First use LinearRegression to train a model with `elastic net` regularization and an intercept. This method returns a `LinearRegressionModel`. Next, use the model's `coefficients` (weights) and `intercept` attributes to print out the model's parameters.

Important! Spark ML (`pyspark.ml`) expects the new vector type `pyspark.ml.linalg.VectorUDT`, but our DataFrame's features column is still using the **old MLLib** type: `pyspark.mllib.linalg.VectorUDT`. So, you need to convert your `parsed_train_data_df` dataframe to be fully compatible with the new Spark ML API before training our linear regression model.

Important! The same conversion will be required in (4c) for the validation set, and also for the training and test sets in (5b), (5c), and (5d) after adding interactions between features

```
In [73]: from pyspark.ml.regression import LinearRegression
from pyspark.mllib.util import MLUtils
# Values to use when training the linear regression model

num_iters = 500 # iterations
reg = 1e-1 # regParam
alpha = .2 # elasticNetParam
use_intercept = True # intercept
```

```
In [75]: # TODO: Replace <FILL IN> with appropriate code

# Convert all vector columns to ML format
parsed_train_data_df = MLUtils.convertVectorColumnsToML(parsed_train_data)

lin_reg = LinearRegression(maxIter=num_iters, regParam=reg, elasticNetParam=alpha)

first_model = lin_reg.fit(parsed_train_data_df)

# coeffsLR1 stores the model coefficients; interceptLR1 stores the model intercept
coeffs_LR1 = first_model.coefficients
intercept_LR1 = first_model.intercept

print(coeffs_LR1, intercept_LR1)
```

```
[23.1982638028573, 26.98504523130359, -67.73180097557051, 52.581253053537, -1
2.602263525137092, -48.387783459855676, 36.41656162742634, -23.60729239535593
4, 0.0, -3.010859551153598, -10.606012321267242, -14.977614644741953] 66.92674
87588785
```

```
In [76]: # TEST LinearRegression (4a)
expected_intercept = 66.92674875888055
expected_weights = [23.198263802857262, 26.985045231303086, -67.73180097557051, -12.602263525139243, -48.387783459857786, 36.41656162742634, 0.0, -3.0108595511539082, -10.606012321268773, -14.977614644741953]
testmti850.Test.assertTrue(np.allclose(intercept_LR1, expected_intercept))
testmti850.Test.assertTrue(np.allclose(coeffs_LR1, expected_weights), 'in
```

```
1 test passed.
1 test passed.
```

(4b) Transform

Now use the `LinearRegressionModel.transform()` method to make predictions on the `parsed_train_data_df`.

```
In [77]: # TODO: Replace <FILL IN> with appropriate code
```

```
sample_prediction = first_model.transform(parsed_train_data_df)

sample_prediction.show()
```

label	features	prediction
0.0	[0.581044957651, 0....]	21.084909469789203
0.0	[0.63944087626, 0....]	63.08554683101235
0.0	[0.665425875415, 0....]	59.18426472167812
0.0	[0.730673247594, 0....]	67.41750658482327
0.0	[0.791114423812, 0....]	74.75083576946234
2.0	[0.56977237281, 0....]	14.55174964241526
2.0	[0.595980071799, 0....]	13.269272145348076
2.0	[0.615847938199, 0....]	38.34650636878545
3.0	[0.424574534176, 0....]	-3.291412687469176
3.0	[0.444046053737, 0....]	7.023831144339326
3.0	[0.483297880403, 0....]	12.67398590639872
3.0	[0.54561164828, 0....]	12.865440995754824
3.0	[0.552845959163, 0....]	6.905799681612088
3.0	[0.601494345313, 0....]	10.475204231854775
4.0	[0.187672926543, 0....]	7.790586512761337
4.0	[0.359595235695, 0....]	18.18359651791863
4.0	[0.400744425698, 0....]	16.054230711571684
4.0	[0.404665218137, 0....]	9.184241129714351
4.0	[0.421594084425, 0....]	30.43956376158438
4.0	[0.423538293212, 0....]	20.915148173273508

only showing top 20 rows

```
In [78]: # TEST Predict (4b)
```

```
testmti850.Test.assertTrue(np.allclose(sample_prediction.first().prediction,
                                         'incorrect value for sample_prediction'))
```

```
1 test passed.
```

(4c) Evaluate RMSE

Next evaluate the accuracy of this model on the validation set. Use the `transform()` method to create predictions, and then use the `calc_RMSE()` function from Part (2b).

```
In [81]: # TODO: Replace <FILL IN> with appropriate code
```

```
parsed_val_data_df = MLUtils.convertVectorColumnsToML(parsed_val_data_df)
val_pred_df = first_model.transform(val_pred_df)

rmse_val_LR1 = calc_RMSE(val_pred_df)
```

```
print ('Validation RMSE:\n\tBaseline = {:.3f} \n\tLR0 = {:.3f}\n\tLR1 =\n\t\t.format(rmse_val_base, rmse_val_LR0, rmse_val_LR1))
```

25/11/28 03:28:22 WARN MLUtils: Vector column conversion has serialization overhead. Please migrate your datasets and workflows to use the spark.ml package.

Validation RMSE:

```
Baseline = 20.759
LR0 = 18.371
LR1 = 16.1152
```

In [82]: # TEST Evaluate RMSE (4c)
`testmti850.Test.assertTrue(np.allclose(rmse_val_LR1, 16.115204285984859),
1 test passed.`

(4d) Grid search

We're already outperforming the baseline on the validation set by almost 2 years on average, but let's see if we can do better. Perform grid search to find a good regularization parameter. Try `regParam` values `1e-10`, `1e-5`, and `1.0`.

In [83]: # TODO: Replace <FILL IN> with appropriate code
`best_RMSE = rmse_val_LR1
best_reg_param = reg
best_model = first_model`
`num_iters = 500 # iterations
alpha = .2 # elasticNetParam
use_intercept = True # intercept`
`regParams = [1e-10, 1e-5, 1]`
`for reg in regParams:
lin_reg = LinearRegression(maxIter=num_iters, regParam=reg, elasticNetParam=alpha, useIntercept=use_intercept)
model = lin_reg.fit(parsed_train_data_df)
val_pred_df = model.transform(parsed_val_data_df)`
`rmse_val_grid = calc_RMSE(val_pred_df)
print(rmse_val_grid)`
`if rmse_val_grid < best_RMSE:
best_RMSE = rmse_val_grid
best_reg_param = reg
best_model = model`
`rmse_val_LR_grid = best_RMSE`

16.122518057355745
16.12251713674809
16.097420376208206

In [84]: # TEST Grid search (4d)
`testmti850.Test.assertTrue(np.allclose(16.097420376208206, rmse_val_LR_gr`
1 test passed.

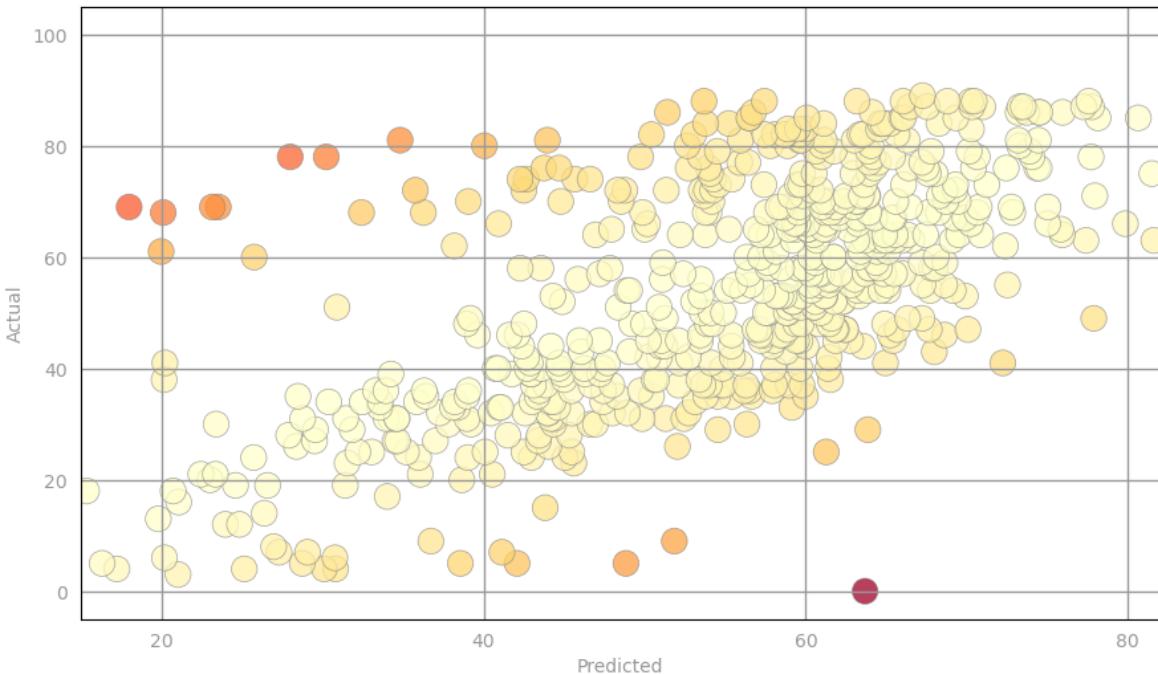
Visualization 5: Best model's predictions

Next, we create a visualization similar to 'Visualization 3: Predicted vs. actual' from Part 2 using the predictions from the best model from Part (4d) on the validation dataset. Specifically, we create a color-coded scatter plot visualizing tuples storing i) the predicted value from this model and ii) true label.

```
In [85]: parsed_val_df = best_model.transform(parsed_val_data_df)
predictions = np.asarray(parsed_val_df
                         .select('prediction')
                         .collect())
actual = np.asarray(parsed_val_df
                     .select('label')
                     .collect())
error = np.asarray(parsed_val_df
                    .rdd
                    .map(lambda lp: squared_error(lp.label, lp.prediction))
                    .collect())

norm = Normalize()
clrs = cmap(np.asarray(norm(error)))[ :, 0:3]

fig, ax = prepareSubplot(np.arange(0, 120, 20), np.arange(0, 120, 20))
ax.set_xlim(15, 82), ax.set_ylim(-5, 105)
ax.scatter(predictions, actual, s=14**2, c=clrs, edgecolors='#888888', alpha=0.8)
ax.set_xlabel('Predicted')
ax.set_ylabel('Actual')
plt.show()
```



Visualization 6: Hyperparameter heat map

Next, we perform a visualization of hyperparameter search using a larger set of hyperparameters (with precomputed results). Specifically, we create a heat map where the brighter colors correspond to lower RMSE values. The first plot has a large area with brighter colors. In order to differentiate within the bright region, we generate a second plot corresponding to the hyperparameters found within that region.

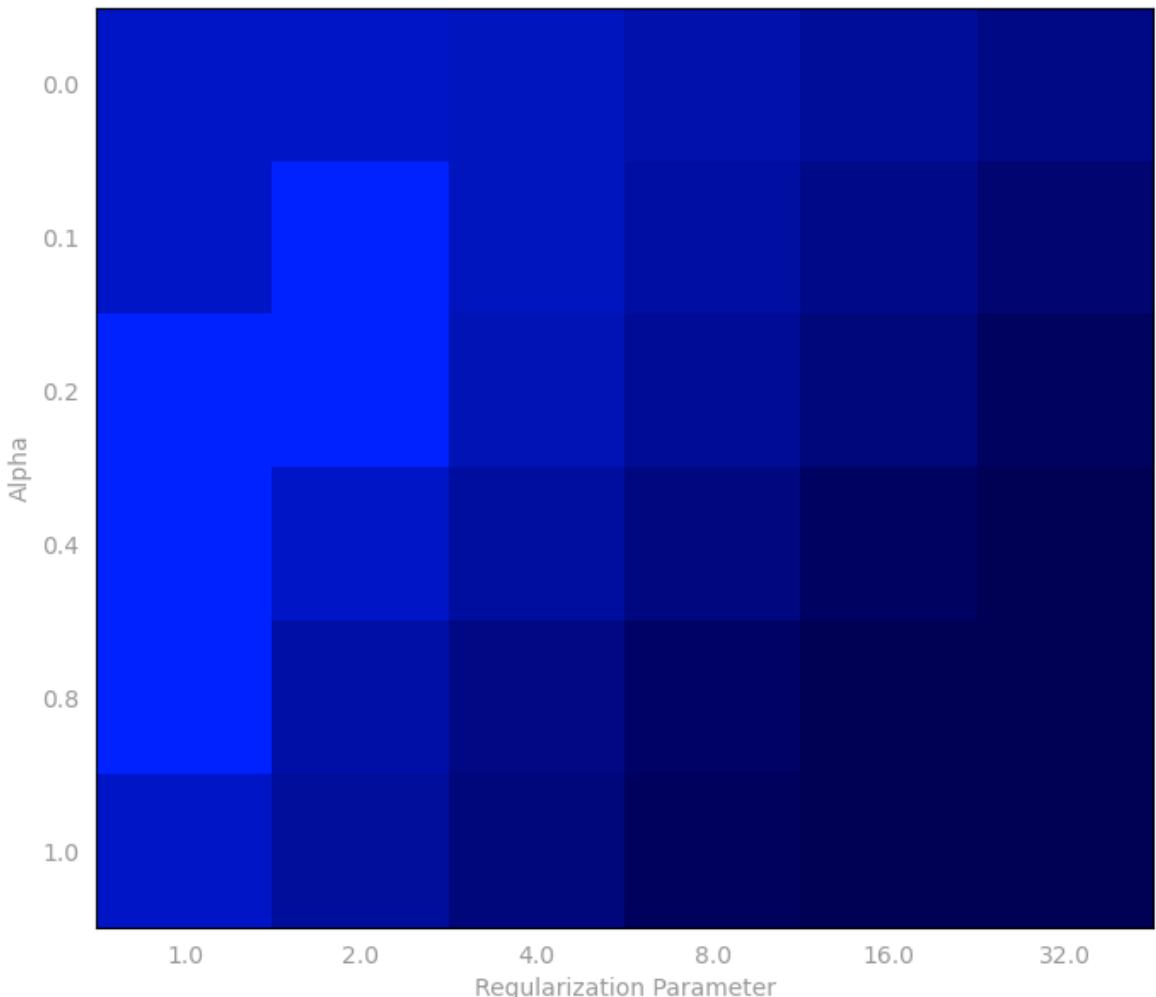
```
In [86]: from matplotlib.colors import LinearSegmentedColormap

# Saved parameters and results, to save the time required to run 36 model
num_iters = 500
reg_params = [1.0, 2.0, 4.0, 8.0, 16.0, 32.0]
alpha_params = [0.0, .1, .2, .4, .8, 1.0]
rmse_val = np.array([[ 15.317156766552452, 15.327211561989827, 15.3571529
                      16.36423857334287, 15.315019185101972, 15.30594921
                      16.231992712117222, 17.700179790697746, 15.3052663
                      15.824676190630191, 17.045905140628836, 19.3655583
                      15.620051033979871, 16.631757941340428, 18.9487868
                      15.522394576046239, 16.414106221093316, 18.6559787
                      15.33442896030322, 15.680134490745722, 16.86502909
                      20.91796910560631 ]])

num_rows, num_cols = len(alpha_params), len(reg_params)
rmse_val = np.array(rmse_val)
rmse_val.shape = (num_rows, num_cols)

fig, ax = prepareSubplot(np.arange(0, num_cols, 1), np.arange(0, num_rows,
                                                               gridWidth=0.))
ax.set_xticklabels(reg_params), ax.set_yticklabels(alpha_params)
ax.set_xlabel('Regularization Parameter')
ax.set_ylabel('Alpha')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055'])
image = ax.imshow(rmse_val, interpolation='nearest', aspect='auto', cmap=colors)
plt.show()
```

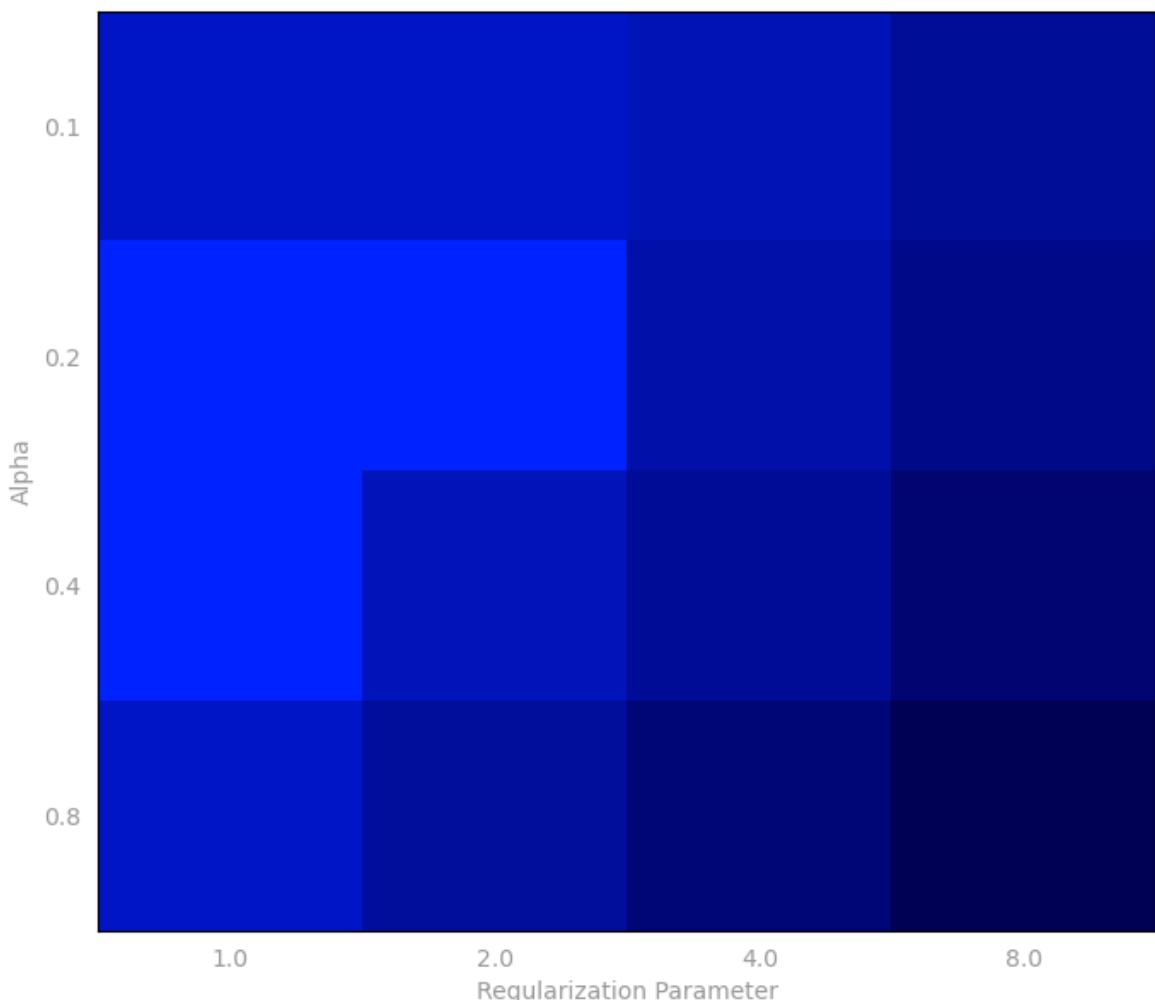


```
In [88]: # Zoom into the top left
alpha_params_zoom, reg_params_zoom = alpha_params[1:5], reg_params[:4]
rmse_val_zoom = rmse_val[1:5, :4]

num_rows, num_cols = len(alpha_params_zoom), len(reg_params_zoom)

fig, ax = prepareSubplot(np.arange(0, num_cols, 1), np.arange(0, num_rows,
                                                               gridWidth=0.))
ax.set_xticklabels(reg_params_zoom), ax.set_yticklabels(alpha_params_zoom)
ax.set_xlabel('Regularization Parameter'), ax.set_ylabel('Alpha')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055'])
image = plt.imshow(rmse_val_zoom, interpolation='nearest', aspect='auto',
                   plt.show())
```



Part 5: Add interactions between features

(5a) Add 2-way interactions

So far, we've used the features as they were provided. Now, we will add features that capture the two-way interactions between our existing features. Write a function `two_way_interactions` that takes in a `LabeledPoint` and generates a new `LabeledPoint` that contains the old features and the two-way interactions between them.

Note:

- A dataset with three features would have nine (3^2) two-way interactions.
- You might want to use `itertools.product` to generate tuples for each of the possible 2-way interactions.
- Remember that you can combine two `DenseVector` or `ndarray` objects using `np.hstack`.

```
In [94]: # TODO: Replace <FILL IN> with appropriate code
import itertools
```

```
def two_way_interactions(lp):
    """Creates a new `LabeledPoint` that includes two-way interactions.
```

Note:

For features $[x, y]$ the two-way interactions would be $[x^2, x*y]$, would be appended to the original $[x, y]$ feature list.

Args:

`lp` (`LabeledPoint`): The label and features for this observation.

Returns:

`LabeledPoint`: The new `LabeledPoint` should have the same label and should include the features from `'lp'` followed by the two-way

"""

`features = lp.features`

```
interaction_features = [x * y for x, y in itertools.product(features,
new_features = np.hstack((features, interaction_features)))
```

```
return LabeledPoint(lp.label, new_features)
```

```
#print(two_way_interactions((0.0, Vectors.dense([2, 3]))))
```

```
# Transform the existing train, validation, and test sets to include two-
# Remember to convert them back to DataFrames at the end.
```

```
train_data_interact_df = parsed_train_data_df.rdd.map(two_way_interactions)
val_data_interact_df = parsed_val_data_df.rdd.map(two_way_interactions).t
test_data_interact_df = parsed_test_data_df.rdd.map(two_way_interactions)
```

```
In [95]: # TEST Add two-way interactions (5a)
```

```
two_way_example = two_way_interactions((0.0, Vectors.dense([2, 3])))
testmti850.Test.assertTrue(np.allclose(sorted(two_way_example[1]),
                                         sorted([2.0, 3.0, 4.0, 6.0, 6.0, 9.0])),
                           'incorrect features generatedBy two_way_interactions')
two_way_point = two_way_interactions((1.0, Vectors.dense([1, 2, 3])))
testmti850.Test.assertTrue(np.allclose(sorted(two_way_point[1]),
                                         sorted([1.0, 2.0, 3.0, 1.0, 2.0, 3.0, 2.0, 4.0, 6.0, 3
                                         ])),
                           'incorrect features generated by twoWayInteractions')
testmti850.Test.assertEquals(two_way_point[0], 1.0, 'incorrect label gene
testmti850.Test.assertTrue(np.allclose(sum(train_data_interact_df.first()
```

```
'incorrect features in train_data_interact_df')

testmti850.Test.assertTrue(np.allclose(sum(val_data_interact_df.first().f
    'incorrect features in val_data_interact_df')

testmti850.Test.assertTrue(np.allclose(sum(test_data_interact_df.first().f
    'incorrect features in test_data_interact_df')
```

-

```
AttributeError Traceback (most recent call last)
Cell In[95], line 2
      1 # TEST Add two-way interactions (5a)
----> 2 two_way_example = two_way_interactions((0.0, Vectors.dense([2,
3])))
      3 testmti850.Test.assertTrue(np.allclose(sorted(two_way_example[1]),
      4                                         sorted([2.0, 3.0, 4.0, 6.0,
9.0])),
      5                                         'incorrect features generatedBy two_way_interactio
ns')
      6 two_way_point = two_way_interactions((1.0, Vectors.dense([1, 2,
3])))

Cell In[94], line 18, in two_way_interactions(lp)
      4 def two_way_interactions(lp):
      5     """Creates a new `LabeledPoint` that includes two-way interact
ions.
      6
      7     Note:
      (...)

      16         should include the features from `lp` followed by the
two-way interaction features.
      17     """
----> 18     features = lp.features
      19     interaction_features = [x * y for x, y in itertools.product(fe
atures, features)]
      20     new_features = np.hstack((features, interaction_features))
```

AttributeError: 'tuple' object has no attribute 'features'

The function 'two_way_interactions' is supposed to take a LabeledPoint as Args, then I commented the test line with only a tuple and I didn't take into account the test code also with tuples.

(5b) Build interaction model

Now, let's build the new model. We've done this several times now. To implement this for the new features, we need to change a few variable names.

Note:

- Remember that we should build our model from the training data and evaluate it on the validation data.
- You should re-run your hyperparameter search after changing features, as using the best hyperparameters from your prior model

will not necessarily lead to the best model.

- For this exercise, we have already preset the hyperparameters to reasonable values.

```
In [98]: # TODO: Replace <FILL IN> with appropriate code
num_iters = 500
reg = 1e-10
alpha = .2
use_intercept = True

train_data_interact_df = MLUtils.convertVectorColumnsToML(train_data_interact)
val_data_interact_df = MLUtils.convertVectorColumnsToML(val_data_interact)

lin_reg = LinearRegression(maxIter=num_iters, regParam=reg, elasticNetParam=alpha, standardization=True, featuresCol="features", labelCol="label", intercept=use_intercept)

model_interact = lin_reg.fit(train_data_interact_df)

preds_and_labels_interact_df = model_interact.transform(val_data_interact_df)

rmse_val_interact = calc_RMSE(preds_and_labels_interact_df)

print ('Validation RMSE:\n\tBaseline = {0:.3f} \n\tLR0 = {1:.3f}\n\tLR1 = {2:.3f}\n\tLRGrid = {3:.3f}\n\tInteract = {4:.3f}'.format(rmse_val_LR_grid, rmse_val_LR0, rmse_val_LR1, rmse_val_LRGrid, rmse_val_Interact))
```

```
Exception ignored in: <_io.BufferedReader name=5>
Traceback (most recent call last):
  File "/opt/spark/python/lib/pyspark.zip/pyspark/daemon.py", line 200, in manager
    BrokenPipeError: [Errno 32] Broken pipe

Validation RMSE:
    Baseline = 20.759
    LR0 = 18.371
    LR1 = 16.115
    LRGrid = 16.097
    Interact = 15.456
```

```
In [99]: # TEST Build interaction model (5b)
testmti850.Test.assertTrue(np.allclose(rmse_val_interact, 15.456123194017))

1 test passed.
```

(5c) Evaluate interaction model on test data

Our next step is to evaluate the new model on the test dataset. Note that we haven't used the test set to evaluate any of our models. Because of this, our evaluation provides us with an unbiased estimate for how our model will perform on new data. If we had changed our model based on viewing its performance on the test set, our estimate of RMSE would likely be overly optimistic.

We'll also print the RMSE for both the baseline model and our new model. With this information, we can see how much better our model performs than the baseline model.

```
In [100...]: # TODO: Replace <FILL IN> with appropriate code
preds_and_labels_test_df = model_interact.transform(test_data_interact_df)
```

```

rmse_test_interact = calc_RMSE(preds_and_labels_test_df)

print ('Test RMSE:\n\tBaseline = {0:.3f}\n\tLRIteract = {1:.3f}'
      .format(rmse_test_base, rmse_test_interact))

```

Test RMSE:

```

Baseline = 20.547
LRIteract = 14.701

```

```
In [101...]: # TEST Evaluate interaction model on test data (5c)
testmti850.Test.assertTrue(np.allclose(rmse_test_interact, 14.70091399650
                                         'incorrect value for rmse_test_interact'))
```

1 test passed.

(5d) Use a pipeline to create the interaction model

Our final step is to create the interaction model using a `Pipeline`. Note that Spark contains the `PolynomialExpansion` transformer which will automatically generate interactions for us. In this section, you'll need to generate the

`PolynomialExpansion` transformer and set the stages for the `Pipeline` estimator. Make sure to use a degree of 2 for `PolynomialExpansion`, set the input column appropriately, and set the output column to "polyFeatures". The pipeline should contain two stages: the polynomial expansion and the linear regression.

```

In [103...]: # TODO: Replace <FILL IN> with appropriate code
from pyspark.ml import Pipeline
from pyspark.ml.feature import PolynomialExpansion

num_iters = 500
reg = 1e-10
alpha = .2
use_intercept = True

polynomial_expansion = PolynomialExpansion(degree=2, inputCol='features',
                                             outputCol='polyFeatures')
linear_regression = LinearRegression(maxIter=num_iters, regParam=reg, elasticNetParam=alpha,
                                      fitIntercept=use_intercept, featuresCol='polyFeatures')

pipeline = Pipeline(stages=[polynomial_expansion, linear_regression])

pipeline_model = pipeline.fit(parsed_train_data_df)

parsed_test_data_df = MLUtils.convertVectorColumnsToML(parsed_test_data_df)
predictions_df = pipeline_model.transform(parsed_test_data_df)

evaluator = RegressionEvaluator()

rmse_test_pipeline = evaluator.evaluate(predictions_df, {evaluator.metricName: 'rmse'})

print('RMSE for test data set using pipelines: {:.3f}'.format(rmse_test_pipeline))

RMSE for test data set using pipelines: 14.675
25/11/28 04:34:13 WARN MLUtils: Vector column conversion has serialization overhead. Please migrate your datasets and workflows to use the spark.ml package.

```

```
In [104...]: # TEST Use a pipeline to create the interaction model (5d)
testmti850.Test.assertTrue(np.allclose(rmse_test_pipeline, 14.67520761192
                                'incorrect value for rmse_test_pipeline'))
```

1 test passed.

Notebook Ended